

Machine, Data and Learning

POMDP Basics

MDP vs. POMDPs



- **MDP:** Agent's percept in any given state identify the state that it is in, e.g., state (4,3) vs (3,3)
 - Given observations, uniquely determine the state
 - Hence, we will not explicitly consider observations, only states
- **POMDP:** Agent's percepts in any given state **DO NOT** identify the state that it is in, e.g., may be (4,3) or (3,3)
 - Given observations, not uniquely determine the state
 - POMDP: Partially observable MDP for inaccessible environments

POMDP: Partially Observable Markov Decision Process

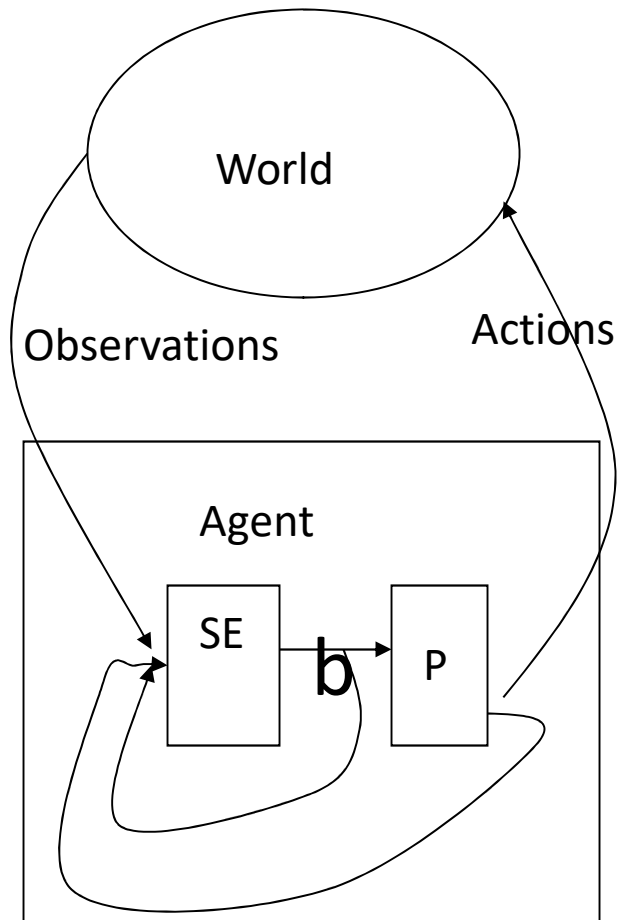
- Set of states, **S**
- Set of actions, **A**
- **P** is the table of transition probabilities
- **R(s,a)** reward received for taking action “a” in state “s”
- **Policy** π maps a **state** “s” to an **action** “a”
- **PLUS**
 - Finite set Ω of observations
 - Table **O** of observation probabilities where **O(o|a,s')** is the probability that “o” is observed given that action “a” taken leads to state s'
 - **Policy** maps **histories of observations** to **actions**

MDP vs POMDP

MDP



POMDP



SE: State estimator

b: Belief state

SE updates the beliefs
based on last observation,
previous belief state
and previous action

P: Policy is no longer a function of the state,
But of the agent's belief state

POMDP:<S, A, P, R, Ω , O>

- **S**, Set of states
- **A**, finite set of actions
- **P** is the table of transition probabilities
- **R(s,a)** reward received for taking action “a” in state “s”
- Finite set Ω of observations, e.g., {red, green} in example below
 - Observations hint at state, e.g., *Observe Red room, but not S3*
- Table **O** of observation probabilities
 - **O(o|a,s')** prob “o” observed given action “a” leads to state s'
 - $P(\text{red} \mid \text{LEFT}, S3) = 0.4$

S1	S2	S3	S4
Green	Red	Red/Green	Green

POMDP: Partially Observable Markov Decision Process



- Agent has initial beliefs
- Agent takes an action
- Gets an observation
- Interprets the observation
- Updates beliefs
- Decides on an action
- Repeats

Agent takes optimal action considering world/other agents

Elements: {States, Actions, Transitions, Rewards, Observations }

POMDP: Partially Observable Markov Decision Process

- Underlying dynamics are still **Markovian**: World has **NOT** changed its characteristics, agents sensors have changed
- Observations only hint at what state we are in, but not exactly identify state
- So, somehow agent may need to remember what it observed in the past and what action it took:
 - *If I observed feature “green” in the past, then took action “left” and then observed “red”, it must mean that I am either in state S_3 (probability of 0.9) or S_2 (Prob 0.1) now*
- *Need to maintain beliefs*

Tiger Problem



- Standing in front of two closed doors
- World is in one of two states: tiger is behind left door or right door
- Three actions: Open left door, open right door, listen
 - *Listening is not free, and not accurate (may get wrong info)*
- Reward: Open the wrong door and get eaten by the tiger (large -ve)
Open the right door and get a prize (small +ve)

Tiger Problem: POMDP Formulation

- Two states: SL and SR
- Three actions: LEFT, RIGHT, LISTEN
- Transition probabilities:

Listen	SL	SR
SL	1.0	0.0
SR	0.0	1.0

Left	SL	SR
SL	0.5	0.5
SR	0.5	0.5

Right	SL	SR
SL	0.5	0.5
SR	0.5	0.5

Tiger Problem: POMDP formulation

- Observations: TL (tiger left) or TR (tiger right)
- Observation probabilities:

Listen	TL	TR
SL	0.85	0.15
SR	0.15	0.85

Left	TL	TR
SL	0.5	0.5
SR	0.5	0.5

Right	TL	TR
SL	0.5	0.5
SR	0.5	0.5

- Rewards:

- $R(SL, Listen) = R(SR, Listen) = -1$
- $R(SL, Left) = R(SR, Right) = -100$
- $R(SL, Right) = R(SR, Left) = +10$

How to Find the Optimal Policy?

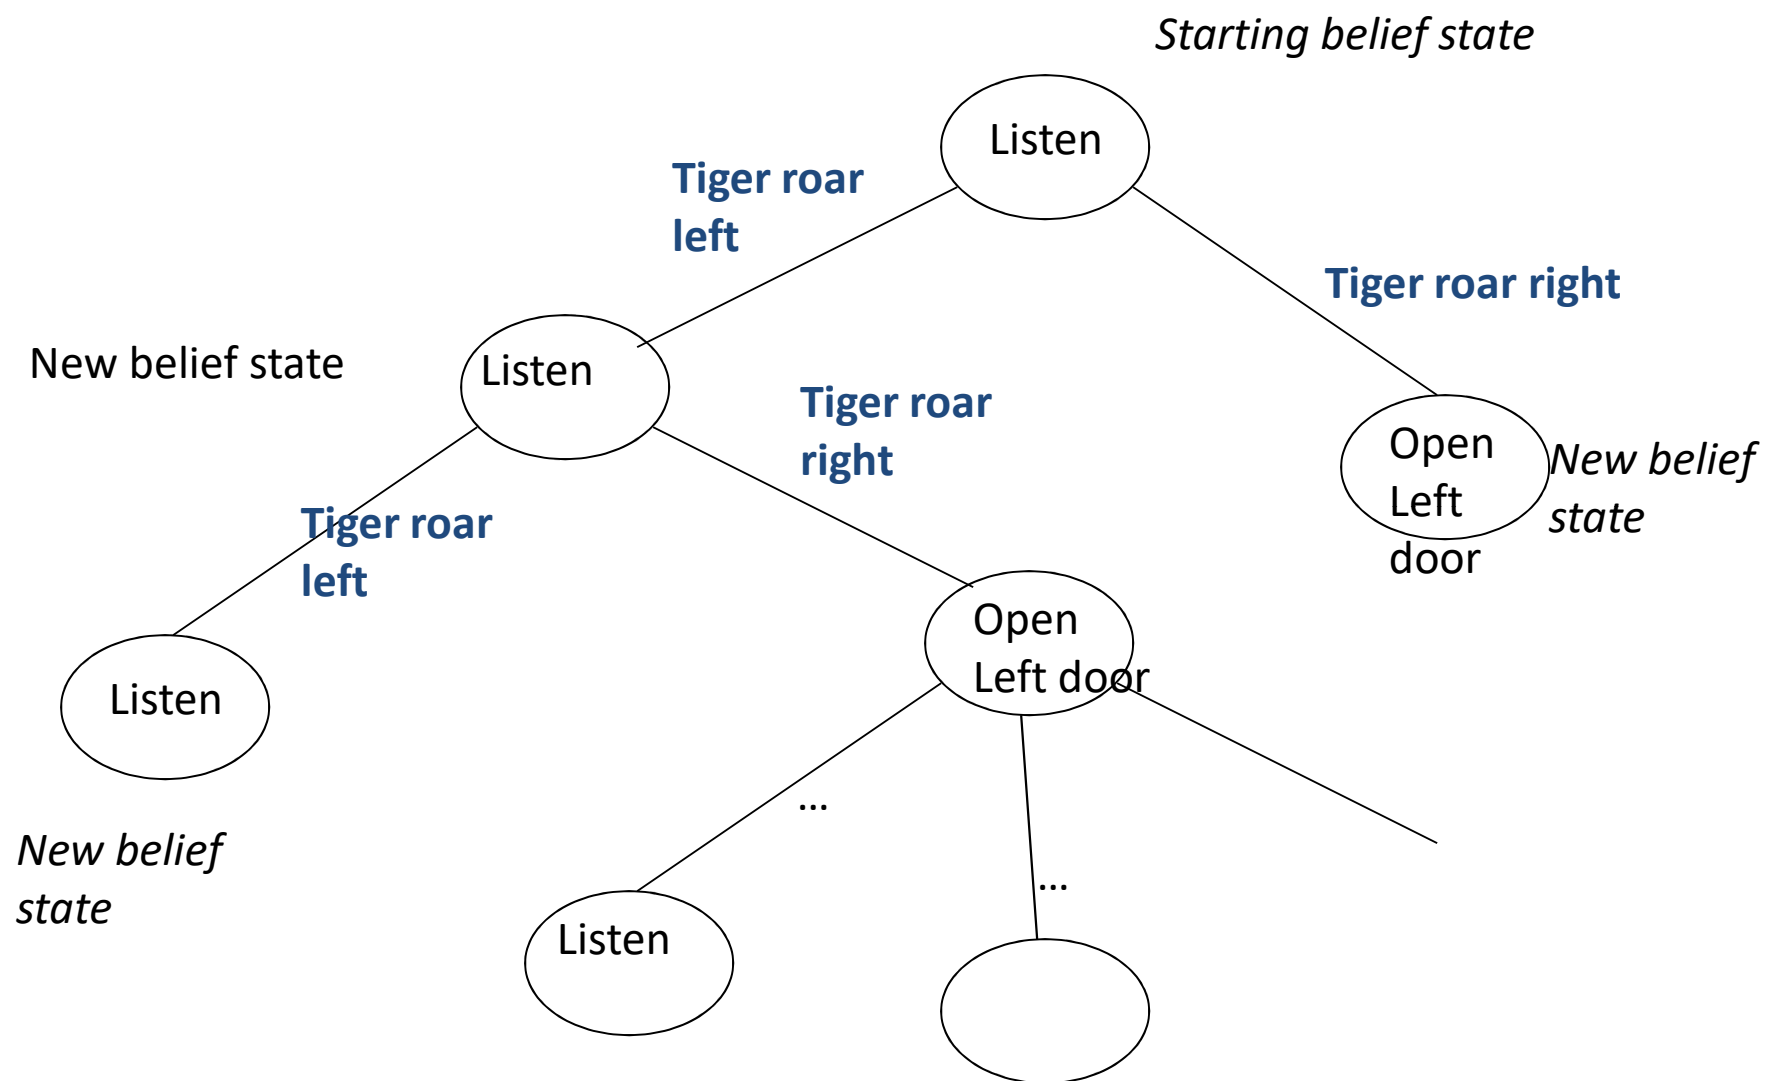
- Now lets find an optimal policy for this problem

- Why not use value iteration directly?

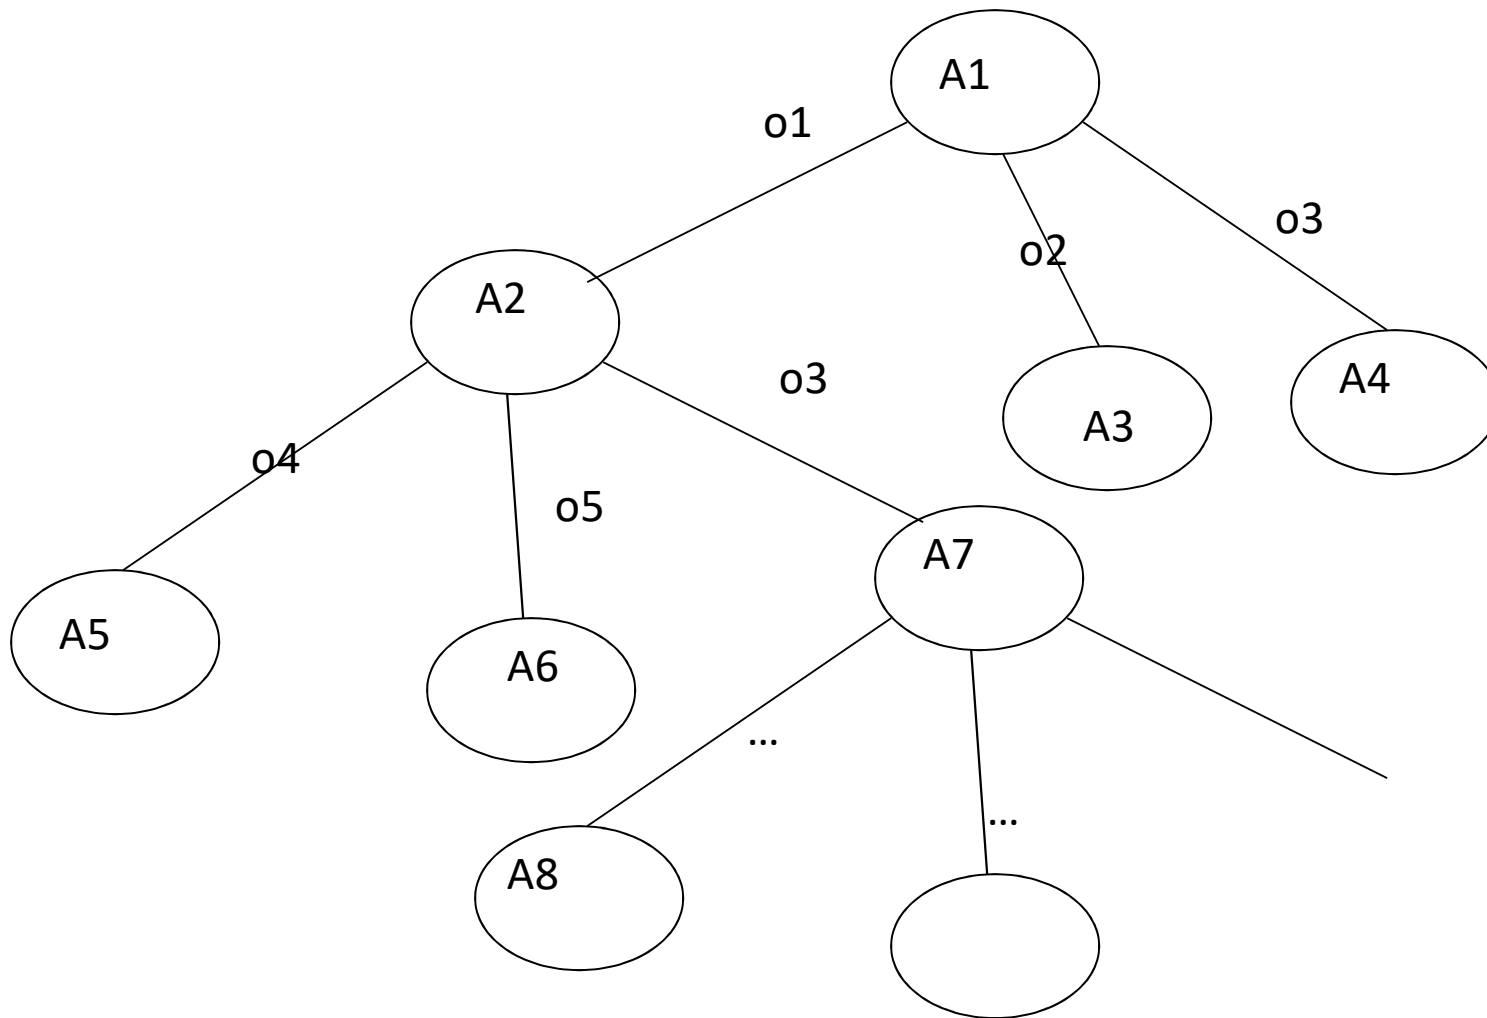
$$- U_{t+1}(I) = R(I) + \max_A \sum_J P(J|I,A) * U_t(J)$$

- Could we compute the utilities in this manner?
- Could such utilities be actually used?
- Need mapping from belief states to actions!

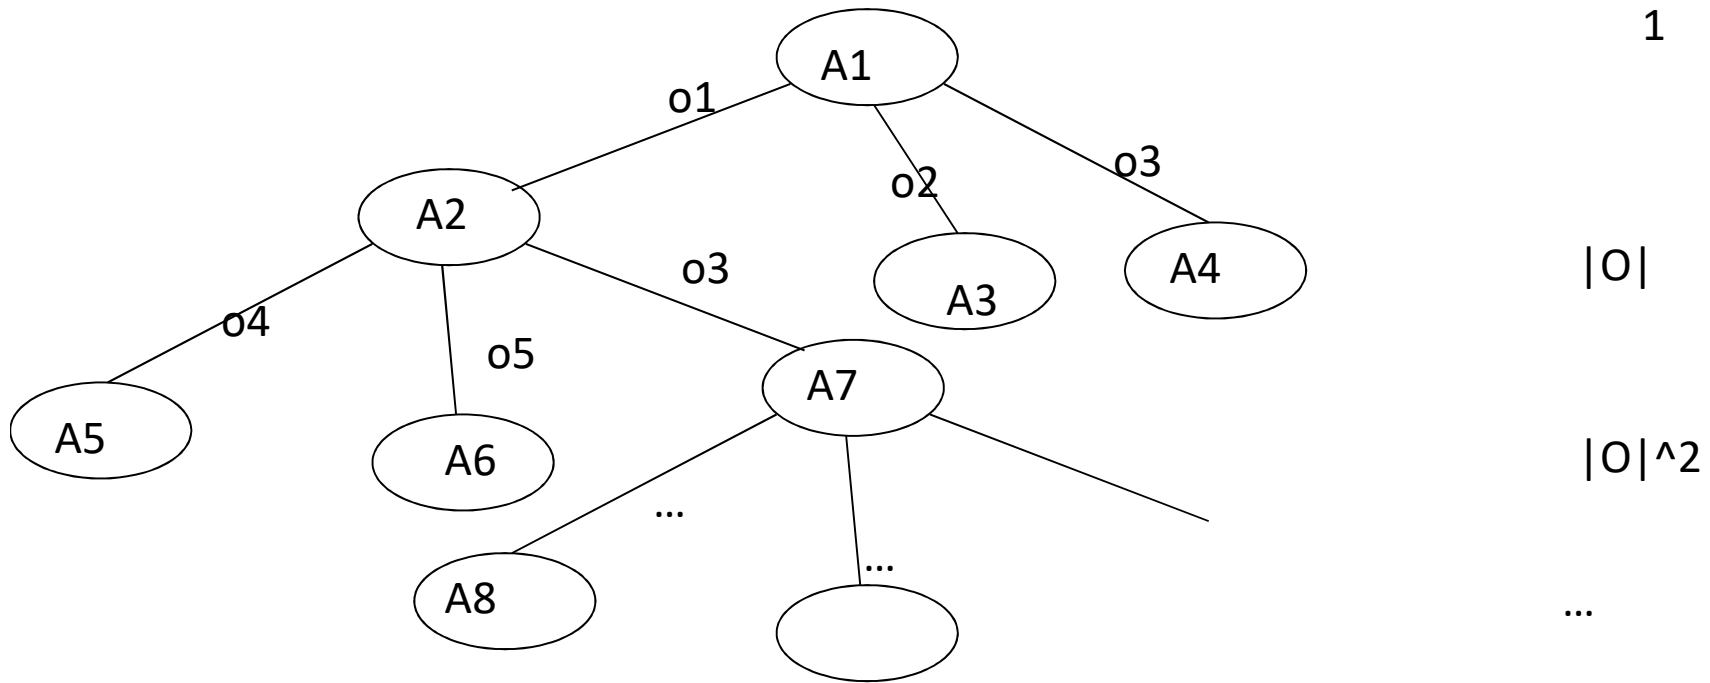
Sample POMDP Policy Tree



POMDP Policy Tree



How Many POMDP policies possible



How many policy trees, if $|A|$ actions, $|O|$ observations, T horizon:

- How many nodes in a tree:

How many trees:


$$N = \sum_{i=0}^{T-1} |O|^i = \frac{(|O|^T - 1)}{(|O| - 1)}$$

$$|A|^N$$

POMDP Belief State

- Computing belief state important, since policy maps belief state to action
 - *Not just the most probable state of the world*
- Probability distributions over the states of the world
 - *Sufficient statistic* for the past history and initial belief state:
 - *No additional data about past actions & observations supplies any further information*
 - *That is, process over belief states is a markov process (why?)*
 - *As if maintained a complete history of actions & observations*

Evolution of Belief State: 1



- Set of states, S_1, S_2, S_3, S_4
- For each $s \in S$, A_s set of actions: Down or Up
- Transition Prob T : 0.9 (direction of move), 0.1 opp
- $R(s,a)$ reward received

– Finite set Ω of observations

– O observation probabilities:

$$\text{– } \Pr(o_1 | s_1) = \Pr(o_1 | s_2) = \Pr(o_1 | s_4) = 1$$

$$\text{– } \Pr(o_2 | s_3) = 1$$

Initial belief state: $[0.333, 0.333, 0, 0.333]$


$S_1 \quad S_2 \quad S_3 \quad S_4$

Evolution of Belief State: 2

Suppose agent moves down and observes o1:

- What should the agent believe about its state? Does it now know more about where it is more likely to be?

- [0.100, 0.450, 0, 0.450]



Belief State

- $b \rightarrow$ probability distribution over our set of states, e.g., over s_1, s_2, s_3, s_4
- $b(s)$ denotes the probability assigned to world state s by belief state b
- *In $[0.333, 0.333, 0, 0.333]$, what is $b(s_1)$?*
- $0 \leq b(s) \leq 1$
$$\sum b(s) = 1$$

$$s \in S$$

Computing Belief States

- s = old state
- b = old belief state, and $b(s)$ probability of s given belief state b
- a = action
- b' = new belief state
- $b'(s')$ = probability of s' given b'
- o = observation

Computing Belief States

$$\begin{aligned} b'(s') &= \Pr(s' \mid o, a, b) = \Pr(s' \wedge o \wedge a \wedge b) / \Pr(o \wedge a \wedge b) \\ &= \frac{\Pr(o \mid s', a, b) \Pr(s' \mid a, b) * \Pr(a \wedge b)}{\Pr(o \mid a, b) * \Pr(a \wedge b)} \\ &= \frac{\Pr(o \mid s', a) \Pr(s' \mid a, b)}{\Pr(o \mid a, b)} \end{aligned}$$

Will not repeat $\Pr(o \mid a, b)$ in the next slide, but it is there!

- *Treated as a normalizing factor, so that b' sums to 1*

Computing Belief States: Numerator

$$= \Pr(o \mid s' a) \Pr(s' \mid a, b) = O(s', a, o) \Pr(s' \mid a, b)$$

$$= O(s', a, o) \sum \Pr(s' \mid a, b, s) \Pr(s \mid a, b)$$

$$= O(s', a, o) \sum \Pr(s' \mid a, b, s) b(s) \quad ; \Pr(s \mid a, b) = \Pr(s \mid b) = b(s)$$

$$= O(s', a, o) \sum T(s, a, s') b(s)$$

(Please work out some of the details later)

Belief State

Overall formula

$$= \frac{O(s', a, o) \sum_s T(s, a, s') b(s)}{\text{Pr}(o \mid a, b)}$$

Example

Moves down and does not observe s_3

- $b \rightarrow b'$
- i.e., $[0.333, 0.333, 0, 0.333] \rightarrow [0.1, 0.45, 0, 0.45]$

$b'(s_1)$ = probability of s_1 in our new belief state b'

Numerator = $\Pr(o_1 \mid s_1, \text{down})$ *

$$[\Pr(s_1 \mid s_1, \text{down}) * b(s_1) + \Pr(s_1 \mid s_2, \text{down}) * b(s_2) + \\ \Pr(s_1 \mid s_3, \text{down}) * b(s_3) + \Pr(s_1 \mid s_4, \text{down}) * b(s_4)]$$

$$= 1 * [0.1 * 0.333 + 0.1 * 0.333 + 0 + 0] = 0.0666$$

Why is this not 0.1?

Example

Moves down and observes o1 (i.e., not observe s3)

- $b \rightarrow b'$
- i.e., $[0.333, 0.333, 0, 0.333] \rightarrow [0.1, 0.45, 0, 0.45]$

In $b'(s1)$, *Numerator* = 0.0666

The above is unnormalized probability, hence not 0.1!

Denote unnormalized $b'(s1)$ as $Ub'(s1)$

- Similarly calculate unnormalized $Ub'(s2)$, $Ub'(s3)$, $Ub'(s4)$
- $[Ub'(s1) + Ub'(s2) + Ub'(s3) + Ub'(s4)]/\text{denominator} = 1$
- Denominator = 0.666 (please check at home)
- $b(s1) = 0.0666/0.666 = 0.1$

Policy: Map Belief State to Action

Convert POMDP \rightarrow “Belief MDP”

- *Recall, process over belief states is markov*
- B, the set of belief states, is the set of MDP states
- A, the set of actions, is the same
- $R'(b,a)$ is the reward function on the belief states:

$$R'(b, a) = \sum_{s \in S} b(s) R(s, a)$$

- Transition function:

$$T(b, a, b') = \Pr(b' \mid a, b) = \sum_{o \in \Omega} \Pr(b' \mid a, b, o) * \Pr(o \mid a, b)$$

$$\begin{aligned} \text{Where } \Pr(b' \mid b, a, o) &= 1 \text{ if } SE(b, a, o) = b' \\ &= 0 \text{ otherwise} \end{aligned}$$

Transition Function

Note: observe-s3 = o2, not(observe-s3) = o1

E.g., $T([0.330, 0.330, 0, 0.330], \text{down}, [0.1, 0.45, 0, 0.45])$
 $= Pr(b' \mid \text{down}, b, \text{observe-s3}) * pr(\text{observe-s3} \mid \text{down}, b)$
 $+ Pr(b' \mid \text{down}, b, \text{Not}(\text{observe-s3})) * pr(\text{Not}(\text{observe-s3}) \mid \text{down}, b)$

$= 0 * pr(\text{observe-s3} \mid \text{down}, b) + 1 * pr(\text{Not}(\text{observe-s3}) \mid \text{down}, b)$

$= Pr(\text{Not}(\text{observe-s3}) \mid \text{down}, b) = 0.666$

$T(b, a, b') = Pr(o \mid a, b)$ where when action “a” taken in belief-state “b” and we observed “o”, we ended up in belief-state b’

Try Value Iteration

- Given belief MDP, if we can generate an optimal policy, it will give rise to optimal behavior for the original POMDP
- How about trying value iteration in this belief MDP?

$$\begin{aligned} V'(b) &= \max [r(b,a) + \sum P(b' \mid b,a) V(b')] \\ &= \max [r(b,a) + \sum P(o \mid b,a) V(b^a_o)] \end{aligned}$$

- *Where $r(b, a) = \sum r(s,a)b(s)$ is the expected immediate reward for taking action a in belief state b*
- *o is the observation, $P(o \mid b, a)$ implies the probability of observing o given action a in belief state b*
- *$V(b^a_o)$ denotes the value for belief state at the next point in time given that action a was taken in belief state b , with observation o*

Problem in Value Iteration

- Infinite possible belief states:
 - Assume: we don't have a fixed start belief state
 - *MDP has a continuous state space*
 - *No longer a table of states where we can maintain a value per state*
- Also, how to back up values of future belief states --- there are too many (infinite) future belief states as well